

ReiOps makes AI agents safe for businesses.

Managed workflows with sandboxes, approval gates, live traces, audit logs, and Telegram operator control.

The product thesis

Agents are becoming abundant. Businesses will pay for the trusted layer that lets them run agents safely.

Live demo: reiops.xyz/demo

GitHub: github.com/typakon4/reiops-demo

Problem: raw agents are hard to trust.

Companies do not just need more agents. They need control, accountability, and support.

- Agents need access to codebases, CRMs, infra, customer data, and internal tools.
- Teams need to know what an agent did, what changed, and what it cost.
- Risky actions need human approval before PRs, deployments, or external messages.
- SMBs and enterprises want reliable outcomes, not unmanaged scripts.

The scarce layer is no longer the agent. The scarce layer is trust.

Solution: managed agent workflows.

ReiOps packages agents into vetted, observable, human-approved business workflows.

1. Vetted workflow

Repeatable workflows for specific business tasks, not generic agent chaos.

2. Control plane

Sandboxing, tool permissions, live trace, costs, and policy gates.

3. Human approval

Operators approve critical actions from the web dashboard or Telegram.

4. Audit trail

Every run becomes reviewable: inputs, actions, approvals, and outcome.

ReiOps = managed marketplace + control plane for business-ready agents.

Demo: GitHub Issue Fixer.

A concrete workflow that shows how ReiOps controls agent execution.

- A GitHub issue is assigned to an AI workflow.
- The agent runs in a sandbox, edits code, and runs tests.
- ReiOps records a live execution trace with files, commands, risk, and cost.
- PR creation is blocked until a human approves it.
- Approval can happen in the web dashboard or Telegram bot.

Live: reiops.xyz/demo • Telegram: [@ReiOps_bot](https://t.me/ReiOps_bot)

Why now: agent capability is crossing into real work.

As agents get tool access, businesses need governance before adoption can scale.

Agent supply is exploding

Open-source and commercial agents are rapidly becoming available to every team.

Tool access creates risk

Agents increasingly touch repos, SaaS apps, databases, tickets, and infrastructure.

Trust becomes the bottleneck

The winning layer is permissions, monitoring, approvals, and support.

We are not betting on one agent. We are building the layer that runs many agents safely.

Business model: start managed, then productize.

A tech-enabled service wedge that becomes a workflow marketplace and control plane.

- SMB packages: monthly managed workflows for common tasks.
- Enterprise: custom deployments with policy, audit, integrations, and SLA.
- Platform: repeatable deployments become workflow templates.
- Marketplace: curated distribution for vetted business-ready agents.

Services teach us what businesses trust. Software scales the repeatable workflows.

Team: technical builder + operations credibility.

A family team combining agent system building with real-world infrastructure operations.

Vsevolod Malakhov — technical founder

18-year-old Information Systems student. Built Rei, a Telegram-first internal agent operating system, and OpenKey, an end-to-end encrypted messenger. Focus: security, distributed systems, agent orchestration, and human-in-the-loop UX.

Vadim Solovev — operations / business lead

15+ years across critical infrastructure, consulting, construction, and data center operations. Brings enterprise process design, budgets, stakeholder management, and operational discipline.

Proof today: live demo, Telegram bot, public repo, deployed web/API/bot stack.

Next: turn the demo into design partner workflows.

The goal is to learn from real teams and productize repeatable safe automation.

- Run controlled GitHub workflows with design partners.
- Add real integrations with scoped permissions and audit export.
- Build policy engine for risky actions and approval roles.
- Expand into a curated marketplace of managed business workflows.

ReiOps is the trusted operations layer for business-ready AI agents.